

Machine Learning with Kernels in Python

J. Andrew Howe, PhD

Experienced Quantitative Modeler
Certified Energy Risk Professional

Mimar Sinan Guzel Sanatlar Universitesi
March 2016

- ▶ Three major topics intertwined in today's lecture
 1. Fundamental machine learning concepts
 2. Unsupervised classification with Kernels
 3. Implementation in Python
- ▶ These will be demonstrated on Fisher's famous Iris dataset

- ▶ *Kernel* refers to a *Reproducing Kernel Hilbert Space* (RKHS)
- ▶ First investigated by the mathematician Aronszajn in 1950.
- ▶ Applying a kernel to some data $X \in \mathbb{R}^{n \times p}$ correspond to nonlinearly mapping X into a higher dimensional *feature space* \mathbb{F} and then taking the dot product in this space
- ▶ Consider the map

$$\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$\phi([x_1, x_2]') = [x_2^2, \sqrt{2}x_1x_2, x_1^2]'$$


- ▶ For two vectors x_i and x_j , we have

$$\begin{aligned} \phi(x_i)' \phi(x_j) &= [x_{i2}^2, \sqrt{2}x_{i1}x_{i2}, x_{i1}^2] [x_{j2}^2, \sqrt{2}x_{j1}x_{j2}, x_{j1}^2]' \\ &= [x_{i2}^2x_{j2}^2, 2x_{i1}x_{i2}x_{j1}x_{j2}, x_{i1}^2x_{j1}^2] \\ &= (x_i'x_j)^2 \end{aligned}$$


- ▶ This is called the *homogenous quadratic kernel*

- ▶ Instead of mapping the data, then computing the dot product, we can compute this function $K(x_i, x_j) = (x_i'x_j)^2$
- ▶ This is called the *kernel trick*, and results in an $n \times n$ -sized *kernel* or *Gram* matrix K .
- ▶ If our original data had $p = 4$ variables and $n = 150$ observations, K will be 150×150 .
- ▶ Further statistical modeling is performed on K instead of X

- ▶ The kernel trick seems intuitively a bad idea
- ▶ As statisticians, we are frequently interested in *dimensional reduction* - *feature selection*, *feature extraction*, *feature engineering*
- ▶ However, the kernel trick inflates the dimensionality of our data
- ▶ Additionally, there are other issues that working with kernels cause:
 - ▶ K scales up rapidly with n (size and computation time)
 - ▶ K is often nonsingular and inversion is numerically unstable
 - ▶ The kernel trick is not reversible
 - ▶ Can't easily predict on new data
- ▶ The benefit we get from this is that clusters / groups in data can often be better separated in the higher-dimensional kernel space
- ▶ This is because the elements of the K matrix are all functions of distances between all pairs of observations



(a) Iris Data



(b) 1st two Variables after Kernel Trick

Figure: Demonstrating Separation in Kernel Space; x is Setosa, o is Versicolor, and * is Virginica

Some Kernel Functions

Function	Form
Linear Polynomial	$(\gamma x_i' x_j + c_0)^1$
Quadratic Polynomial	$(\gamma x_i' x_j + c_0)^2$
Cubic Polynomial	$(\gamma x_i' x_j + c_0)^3$
RBF (Gaussian)	$\exp(-\gamma \ x_i - x_j\ ^2)$
Sigmoid	$\tanh(\gamma x_i' x_j + c_0)$
Laplace	$\exp(-\gamma \ x_i - x_j\ _1)$
Chi ²	$\exp\left(-\gamma \sum_i \frac{(x_i - x_j)^2}{x_i + x_j}\right)$

- ▶ We will perform Kernel Support Vector Machine (SVM) classification on the Iris dataset
- ▶ Along with seeing the performance of Kernel SVM for classification, we will see several general machine learning techniques
- ▶ I have written my own code (mostly in MATLAB) to do everything we'll see here - the kernel trick, supervised classification, cross-validation, grid search, feature selection
- ▶ Knowing how to code for machine learning can be invaluable for all of us
- ▶ However, we will use the *scikit-learn* machine learning package in Python, which does much of it and makes automation of machine learning methods easy

- ▶ First, we use Logistic Regression to build a supervised classification model for the Iris Data, with the result:
Logistic Regression: 96.00%
- ▶ Now, let's use the SVM with the linear polynomial, quadratic polynomial, RBF, and sigmoid kernels to classify the iris data
- ▶ Using the `svm.SVC` object which implements Kernel Support Vector Machine classification, we have:
Linear Kernel: 99.33%
Quadratic Kernel: 98.67%
RBF Kernel: 98.67%
Sigmoid Kernel: 4.00%
- ▶ But there is a possibility that these correct classification rates are all inflated because the model has been *overfit* to the observed data
- ▶ We can fix this with *cross-validation*

- ▶ Cross-validation (CV) refers to a general technique where the data is partitioned into different groups:

Training Data used to fit the model

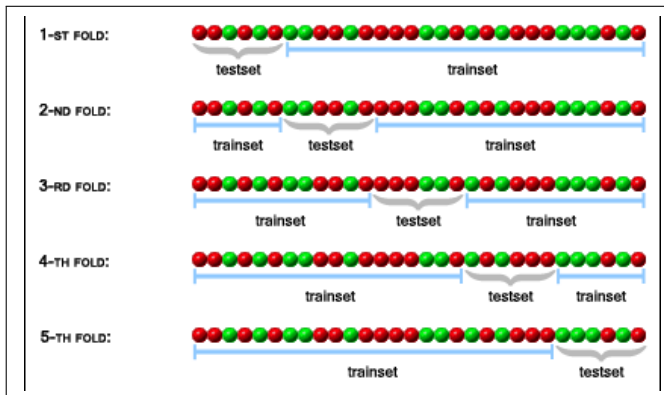
Testing Data used to measure the performance of the model

Validation Data used to validate learning for iterative-algorithms (optional, used less frequently)

- ▶ There are three broad types of cross-validation
 - ▶ leave- p out
 - ▶ K -fold
 - ▶ randomized

- ▶ With leave- p out cross validation, we create $\binom{n}{p}$ partitions of the data
- ▶ For each, the training set has $n - p$ observations, with the remaining p as the testing set
- ▶ For large n and p , it becomes impractical to create all leave- p out CV sets
- ▶ If $n = 150$ and $p = 15$: $1.6239221627803363e+20$ sets

- ▶ in K -fold CV, the data is divided into partitions such that every training set is different



An Example of 5-fold CV

- ▶ The data is often split into K -folds many times, with the first fold being randomly selected each time

- ▶ In randomized CV, a certain percentage of the data is randomly (uniformly) selected for inclusion in the training set
- ▶ The remaining observations are used as the testing set
- ▶ Since every observation has an equal chance of being selected for training, there's no guarantee that every partition will be different
- ▶ I generally use randomized cross-validation
- ▶ The scikit-learn package makes it easy to implement all these (and more) types of cross-validation

- ▶ Using the *cross_validation.ShuffleSplit* object and the *cross_validation.cross_val_score* function, the correct classification for each kernel is:

Correct Classification Summary

	Linear	Quadratic	RBF	Sigmoid
Min.	93.3%	90.0%	91.7%	0.0%
Mean	97.5%	95.1%	96.8%	26.5%
Median	98.3%	96.7%	96.7%	28.3%
Max.	100.0%	100.0%	100.0%	31.6%

- ▶ But what about those parameters we set for the kernel functions - could it be those are affecting the classification performance?
- ▶ We need to tune, or optimize, the parameters to identify a model with the best performance

- ▶ We defined four Kernel SVM objects with static parameters

Some Kernel Functions

Kernel	Parameters
Polynomial	$d = 1, c_0 = 0$
Polynomial	$d = 2, c_0 = 0$
RBF (Gaussian)	$\gamma = \frac{1}{p}$
Sigmoid	$\gamma = \frac{1}{p}, c_0 = 0$

- ▶ Perhaps a non-homogenous cubic polynomial kernel ($d = 3, c_0 \neq 0$) would outperform the linear?
- ▶ What if we set $\gamma = 1$ for the RBF and Sigmoid?
- ▶ There are several different non-stochastic methods to optimize the parameters

- ▶ If we have a countable set of possible reasonable alternative values for each parameter, we could perform a grid search of all combinations of parameters
- ▶ Alternatively, if we know a range of possible values for each parameter, we can iteratively optimize each parameter with a real-valued univariate search against a profile score
- ▶ Yet a third option would be full multivariate optimization
- ▶ We'll see how to easily perform grid search - under cross-validation - to tune the parameters of Kernel SVM for the iris data

- ▶ We want to try all these models:

Polynomial $d = [1, 2, 3]$, $\gamma = [\frac{1}{p}, 1, 2]$, $c_0 = [-1, 0, 1]$

RBF $\gamma = [\frac{1}{p}, 1, 2]$

Sigmoid $\gamma = [\frac{1}{p}, 1, 2]$, $c_0 = [-1, 0, 1]$

- ▶ This is a total of $27 + 3 + 9 = 39$ models, which we evaluate using the *fit* method of the *grid_search.GridSearchCV* object
- ▶ According to this, the model which gives the best correct classification performance uses the non-homogenous Linear Polynomial kernel with $c_0 = -1$
- ▶ Using the same 100 randomized CV partitions, this model obtains an average correct classification of 97.5%, with a median of 98.3%

- ▶ I mentioned that the kernel trick was not reversible
- ▶ After we create a classification model in the feature space \mathbb{F} , there's no way to go back and obtain a classification model in the original data space
- ▶ As well as complicating prediction of new observations, this makes it difficult for us to identify the most influential variables in our dataset, and to perform feature selection
- ▶ To perform feature selection with Kernel SVM classification, we need to sequentially select subsets of features and fit the Kernel model
- ▶ We can do this in Python, simultaneously with cross-validation and parameter tuning, easily

- ▶ Since the Iris data has $p = 4$ variables, there are $2^4 - 1 = 15$ possible subset models, which is low enough to reasonably perform combinatorial subset analysis
- ▶ The scikit-learn package has an object that performs recursive feature extraction, which is like reverse selection for regression
- ▶ It seems that this object won't work with our grid search for parameter tuning, though
- ▶ It's relatively simple to code our own procedure to perform the cross-validated parameter tuning on all possible subsets for this data

- ▶ The result of performing the cross-validated grid search on all possible subsets of the Iris data is:

Best Subset All Variables

Best Model Non-homogenous linear polynomial with $c_0 = -1$

- ▶ For larger datasets, complete enumeration of all possible subsets of features is not reasonable, so more advanced techniques, such as the genetic algorithm, must be used

- ▶ Several of you may be wondering why I'm showing Python, as opposed to MATLAB or R
- ▶ MATLAB is very expensive, and most companies won't pay for it; there is Octave, a FOSS version of MATLAB, but it is not used so much in the real world
- ▶ There are several reasons to prefer Python over R, R:
 - ▶ has a steep learning curve
 - ▶ is not designed to operate well in a business environment
 - ▶ is maintained by statisticians, not software developers (you wouldn't go to the world's best knee surgeon for brain surgery, would you?)
- ▶ Python provides us several benefits:
 - ▶ heavily used by IT and software developers, so it's easy to implement machine learning models in production environments
 - ▶ developed with a focus on readability and productivity
 - ▶ extremely flexible and extensible
- ▶ I use the Anaconda distribution of Python, which automatically installs many common packages used for scientific computing